

MNISTLogReg

March 28, 2018

```
In [13]: import numpy as np
import pandas as pd
from keras.datasets import mnist
from sklearn.linear_model import LogisticRegression
from sklearn import metrics

import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

1 Logistic regression for MNIST

Logistic regression is a model that has many interpretations. The one I will use is the idea that logistic regression minimizes the cross-entropy between predictions of a linear function passed through a normalizing function and the true probabilities. The cross entropy loss for a single prediction over C categories is as follows:

$$L(y, \hat{y}) = \sum_{k=1}^C y_k \log(\hat{y}_k) \quad (1)$$

If we have more examples, we just take the average over them. A linear regression model will not suffice here, because \hat{y} are predicted probabilities, so they must be normalized. This can be done by using a softmax function. If we say we get some $u = Wx$, where x is an input vector and W maps x to \mathbb{R}^C , then

$$\hat{y}_j = \text{softmax}(u)_j = \frac{e^{u_j}}{\sum_{k=1}^C e^{u_k}} \quad (2)$$

Due to memory constraints, we subsample a set of 10000 examples for training this model. We learn the weight matrix W that minimizes the average cross-entropy loss. We can see that we get an accuracy of about 91 percent. A confusion matrix demonstrating the complete test results can be seen at the bottom.

```
In [2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train.reshape(-1, 28**2) / 255.0
x_test = x_test.reshape(-1, 28**2) / 255.0
```

```

In [3]: # subsample x_train
n_train_per_label = 1000
new_examples = []
for i in range(10):
    subset = x_train[y_train==i]
    choices = np.random.choice(len(subset), n_train_per_label, False)
    new_examples.append(subset[choices])
x_train = np.concatenate(new_examples, 0)
y_train = np.repeat(np.arange(10), n_train_per_label)

In [4]: model = LogisticRegression(multi_class='multinomial', solver='lbfgs', verbose=True, max_iter=1000)

In [5]: model.fit(x_train, y_train)
pred_y = model.predict(x_test)

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 6.9s finished

Out[5]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=500, multi_class='multinomial',
    n_jobs=1, penalty='l2', random_state=None, solver='lbfgs',
    tol=0.0001, verbose=True, warm_start=False)

In [6]: print('accuracy=%.2f'%model.score(x_test, y_test))

accuracy=0.91

In [24]: conf_mat = metrics.confusion_matrix(y_test, pred_y)
class_names = list(range(10))
cm = conf_mat.astype('float') / conf_mat.sum(axis=1)[:, np.newaxis]
df_cm = pd.DataFrame(
    cm, index=class_names, columns=class_names,
)
fig = plt.figure(figsize=(7, 6))
heatmap = sns.heatmap(df_cm, annot=conf_mat, fmt="d")
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right')
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45, ha='right')
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.title('Confusion Matrix')

Out[24]: Text(0.5,1,'Confusion Matrix')

```

